

# 04\_homework\_linear\_regression

February 2, 2018

## 1 Programming assignment 4: Linear regression

```
In [1]: import numpy as np
```

```
from sklearn.datasets import load_boston
from sklearn.model_selection import train_test_split
```

### 1.1 Your task

In this notebook code skeleton for performing linear regression is given. Your task is to complete the functions where required. You are only allowed to use built-in Python functions, as well as any numpy functions. No other libraries / imports are allowed.

### 1.2 Load and preprocess the data

In this assignment we will work with the Boston Housing Dataset. The data consists of 506 samples. Each sample represents a district in the city of Boston and has 13 features, such as crime rate or taxation level. The regression target is the median house price in the given district (in \$1000's).

More details can be found here: <http://lib.stat.cmu.edu/datasets/boston>

### 1.3 Task 1: Fit standard linear regression

```
In [2]: def fit_least_squares(X, y):
        """Fit ordinary least squares model to the data.

        Parameters
        -----
        X : array, shape [N, D]
            (Augmented) feature matrix.
        y : array, shape [N]
            Regression targets.

        Returns
        -----
        w : array, shape [D]
            Optimal regression coefficients (w[0] is the bias term)."""
```

```

"""
X_XTranspose_product=np.dot(np.transpose(X),X);
X_XTranspose_product_as_matrix=np.asmatrix(X_XTranspose_product);
X_XTranspose_product_inverse=X_XTranspose_product_as_matrix.getI();
X_XTranspose_product_inverse_asndarray = np.squeeze(np.asarray(X_XTranspose_product_inverse));
w=np.dot(X_XTranspose_product_inverse_asndarray,np.dot(np.transpose(X),y));
return w

```

## 1.4 Task 2: Fit ridge regression

```

In [3]: def fit_ridge(X, y, reg_strength):
        """Fit ridge regression model to the data.

        Parameters
        -----
        X : array, shape [N, D]
            (Augmented) feature matrix.
        y : array, shape [N]
            Regression targets.
        reg_strength : float
            L2 regularization strength (denoted by lambda in the lecture)

        Returns
        -----
        w : array, shape [D]
            Optimal regression coefficients (w[0] is the bias term).

        """
        X_size=np.shape(X);
        lambda_plus_X_XTranspose=np.add(np.multiply(reg_strength,np.identity(X_size[1])),np.
        lambda_plus_X_XTranspose_as_matrix=np.asmatrix(lambda_plus_X_XTranspose);
        lambda_plus_X_XTranspose_inverse=lambda_plus_X_XTranspose_as_matrix.getI();
        lambda_plus_X_XTranspose_inverse_asndarray = np.squeeze(np.asarray(lambda_plus_X_XTranspose_inverse));
        w=np.dot(lambda_plus_X_XTranspose_inverse_asndarray,np.dot(np.transpose(X),y));
        return w;

```

## 1.5 Task 3: Generate predictions for new data

```

In [4]: def predict_linear_model(X, w):
        """Generate predictions for the given samples.

        Parameters
        -----
        X : array, shape [N, D]
            (Augmented) feature matrix.
        w : array, shape [D]
            Regression coefficients.

```

```

Returns
-----
y_pred : array, shape [N]
    Predicted regression targets for the input data.

"""
y_pred=np.dot(X,w);
return y_pred

```

## 1.6 Task 4: Mean squared error

```

In [5]: def mean_squared_error(y_true, y_pred):
        """Compute mean squared error between true and predicted regression targets.

        Reference: https://en.wikipedia.org/wiki/Mean\_squared\_error`

        Parameters
        -----
        y_true : array
            True regression targets.
        y_pred : array
            Predicted regression targets.

        Returns
        -----
        mse : float
            Mean squared error.

        """
        y_true_size=np.shape(y_true);
        mse=np.sum(np.power(np.subtract(y_pred,y_true),2))/y_true_size[0];
        return mse;

```

## 1.7 Compare the two models

The reference implementation produces \* MSE for Least squares  $\approx 23.98$  \* MSE for Ridge regression  $\approx 21.05$

Your results might be slightly (i.e.  $\pm 1\%$ ) different from the reference solution due to numerical reasons.

```

In [6]: # Load the data
        np.random.seed(1234)
        X , y = load_boston(return_X_y=True)
        X = np.hstack([np.ones([X.shape[0], 1]), X])
        test_size = 0.2
        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=test_size)

        # Ordinary least squares regression

```

```

w_ls = fit_least_squares(X_train, y_train)
y_pred_ls = predict_linear_model(X_test, w_ls)
mse_ls = mean_squared_error(y_test, y_pred_ls)
print('MSE for Least squares = {0}'.format(mse_ls))

# Ridge regression
reg_strength = 1
w_ridge = fit_ridge(X_train, y_train, reg_strength)
y_pred_ridge = predict_linear_model(X_test, w_ridge)
mse_ridge = mean_squared_error(y_test, y_pred_ridge)
print('MSE for Ridge regression = {0}'.format(mse_ridge))

```

MSE for Least squares = 23.984307611777403

MSE for Ridge regression = 21.051487033772723